

## MANAGEMENT SCIENCE

Volume 63 • Number 4 • April 2017



## Management Science

Publication details, including instructions for authors and subscription information:

<http://pubsonline.informs.org>

## Competition-Based Dynamic Pricing in Online Retailing: A Methodology Validated with Field Experiments

Marshall Fisher, Santiago Gallino, Jun Li

To cite this article:

Marshall Fisher, Santiago Gallino, Jun Li (2017) Competition-Based Dynamic Pricing in Online Retailing: A Methodology Validated with Field Experiments. Management Science

Published online in Articles in Advance 27 Jun 2017

<https://doi.org/10.1287/mnsc.2017.2753>

Full terms and conditions of use: <http://pubsonline.informs.org/page/terms-and-conditions>

This article may be used only for the purposes of research, teaching, and/or private study. Commercial use or systematic downloading (by robots or other automatic processes) is prohibited without explicit Publisher approval, unless otherwise noted. For more information, contact [permissions@informs.org](mailto:permissions@informs.org).

The Publisher does not warrant or guarantee the article's accuracy, completeness, merchantability, fitness for a particular purpose, or non-infringement. Descriptions of, or references to, products or publications, or inclusion of an advertisement in this article, neither constitutes nor implies a guarantee, endorsement, or support of claims made of that product, publication, or service.

Copyright © 2017, INFORMS

Please scroll down for article—it is on subsequent pages



INFORMS is the largest professional society in the world for professionals in the fields of operations research, management science, and analytics.

For more information on INFORMS, its publications, membership, or meetings visit <http://www.informs.org>

# Competition-Based Dynamic Pricing in Online Retailing: A Methodology Validated with Field Experiments

Marshall Fisher,<sup>a</sup> Santiago Gallino,<sup>b</sup> Jun Li<sup>c</sup>

<sup>a</sup>The Wharton School, University of Pennsylvania, Philadelphia, Pennsylvania 19104; <sup>b</sup>Tuck School of Business, Dartmouth College, Hanover, New Hampshire 03755; <sup>c</sup>Ross School of Business, University of Michigan, Ann Arbor, Michigan 48109

Contact: fisher@wharton.upenn.edu,  <http://orcid.org/0000-0002-5800-240X> (MF); santiago.gallino@tuck.dartmouth.edu,  <http://orcid.org/0000-0002-8906-2691> (SG); junwli@umich.edu,  <http://orcid.org/0000-0002-9237-9147> (JL)

Received: January 9, 2015

Revised: February 8, 2016; August 1, 2016; November 23, 2016

Accepted: December 19, 2016

Published Online in Articles in Advance: June 27, 2017

<https://doi.org/10.1287/mnsc.2017.2753>

Copyright: © 2017 INFORMS

解决方法

策略试验效果

**Abstract.** A retailer following a competition-based dynamic-pricing strategy tracks competitors' price changes and then must answer the following **questions**: (i) Should we respond? (ii) If so, to whom? (iii) How much of a response? (iv) And on which products? The **answers** require unbiased measures of price elasticity as well as accurate estimates of competitor significance and the extent to which consumers compare prices across retailers. There are two key **challenges** to quantify these factors empirically: first, the endogeneity associated with almost any type of observational data, where prices are correlated with demand shocks observable to pricing managers but not to researchers, and second, the absence of competitor sales information, which prevents efficient estimation of a full consumer-choice model. We **address** the first issue by conducting a field experiment with randomized prices. We **resolve** the second issue by exploiting the retailer's own and competitors' stockouts as a source of variation to the consumer choice set, in addition to variations in competitors' prices. We estimate an empirical **model** capturing consumer choices among substitutable products from multiple retailers. Based on the estimates, we propose and test a best-response pricing strategy through a carefully controlled live experiment that lasts five weeks. The experiment documents an 11% revenue increase while maintaining a margin above a retailer-specified target.

**History:** Accepted by Gad Allon, operations management.

**Funding:** The financial support of the Wharton Global Initiatives program and the Fishman-Davidson Center at the Wharton School of Business is gratefully acknowledged.

**Supplemental Material:** The e-companion is available at <https://doi.org/10.1287/mnsc.2017.2753>.

**Keywords:** online retailing • field experiment • competition-based dynamic pricing • stockouts • consumer choice

## 1. Introduction

The Internet has changed the way price information is disseminated. With just a few clicks, consumers are able to obtain price information from multiple retailers. This increased price transparency induces fierce competition among online retailers and requires real-time monitoring and quick responses to competition (Angwin and Mattioli 2012). Many online retailers in turn adopt a competition-based pricing strategy wherein they constantly monitor competitors' prices and use this information as an input in setting their own prices. For example, they may always charge X dollars or X% less or more than a target competitor or any competitor with the lowest price. Retailers miss several opportunities with such simple heuristics. Instead, they should ask: Shouldn't my reaction depend on consumers' elasticity to prices? Shouldn't my reaction depend on the extent to which consumers compare prices across retailers? Shouldn't my reaction depend on changes in availability at competing retailers? Should I still match prices if it seems like the

competitor made a pricing mistake? We address these questions in this paper.

Determining the best-response price requires knowing how demand reacts to price changes. This is a challenging task. Simply regressing historical sales on prices while controlling for observable product characteristics and seasonality usually suffers from endogeneity issues (see Villas-Boas and Winer 1999). Pricing managers often observe demand signals that we researchers do not, such as unobserved product characteristics or a temporal surge in demand due to manufacturer advertisements, and they may adjust prices based on these signals. If they increase prices when they see a demand surge, this creates a positive correlation that fallaciously implies that a higher price results in higher demand. Moreover, the relationship between demand and price will be further confounded by the price levels of substitutable products that the same retailer offers and the price levels of the same and similar products that competitors offer.

To determine the best-response price, we also need to understand the extent to which consumers compare

需求和价格的关系  
价格弹性

顾客的比价程度

prices across retailers. In a situation where consumers always choose the cheapest retailer for any product they buy, retailers need to either charge the lowest price in the market or accept no sales. However, in a situation where consumers consider both prices and retail channels when making purchase decisions, retailers may not need to match every single price that a competitor charges, especially when considering products with low price elasticities or less relevant competitors. Accurately assessing the level of consumer engagement in price comparison across retailers will enable targeted price responses that are efficient and effective.

We partnered with a leading Chinese online retailer, which we will refer to hereafter as the “retail partner,” to address these challenges. We begin by developing a demand model to understand how consumers make choices when given a set of substitutable products from multiple retailers. Our model resolves a key challenge many retailers face when attempting to implement a choice model to understand consumer purchase decisions: absence of competitor sales information. We exploit the retailer’s own and competitors’ stockouts as a source of variation for identification, in addition to competitive price variations. Because stockouts alter the consumer choice set, they provide additional moment conditions that are instrumental in estimating consumers’ preferences for retailers and their level of engagement in comparing prices across retailers. Next, we conduct a randomized price experiment to obtain unbiased estimates of price elasticities addressing the endogeneity challenge. In this experiment, we randomly assign prices to each product under study using a fractional factorial design. After obtaining model estimates using data generated during the one-month experiment period, we solve a constrained optimization problem to define optimal price responses to competitor price changes. Finally, in collaboration with our retail partner, we evaluate the performance of our best-response pricing algorithm through a carefully controlled field experiment. The daily categorical revenue of the treatment group increases by 11% following our methodology, compared to the control group for the same before and after periods.

### 1.1. Contributions

Our paper contributes to the operations management literature and to retail practice in a number of ways. In what follows, we discuss our contributions and the relevant literature.

**1.1.1. Revenue Management and Dynamic Pricing in Retailing.** Competition-based dynamic pricing is a recent development driven by the competitive nature of online retailing. Conventionally, dynamic pricing has been applied mostly in settings with perishable inventory and a finite selling season. For example, Caro and Gallien (2012) study the development and

清理定价模型

implementation of a clearance pricing model for the fast-fashion retailer Zara. Harsha et al. (2016) propose a new revenue-management model to maximize total profit in omni-channel retailing and, document, on average, 6% improvement in markdown revenue through a proprietary pilot with a major U.S. retailer. Cohen et al. (2017) propose a linear integer programming approximation to solve the dynamic promotional pricing problem, and show that their method can improve profit by 3% using sales data from a grocery retailer.

In our setting, however, the need for dynamic pricing arises not from constrained capacity or a perishable product or service, but from constantly changing market conditions due to competition. This dynamic market environment poses new challenges to, and opportunities for, retail pricing. The longstanding literature of retail pricing focuses mostly on pricing and promotion decisions for a single brick-and-mortar retailer holding competitive prices constant (e.g., Caro and Gallien 2012, Cohen et al. 2017) or a long-term competitive pricing strategy (e.g., Lal and Rao 1997). In these traditional retail settings, the pressure for frequent competitive responses is less prevalent due to the customer’s high search costs and the retailer’s high menu cost of changing prices. Our work expands the existing literature on this topic, since these two factors are not present in an online setting, which makes competition-based dynamic pricing a very relevant issue.

**1.1.2. Consumer-Choice Modeling.** We propose a parsimonious choice model that captures the key tension involved in this competitive environment. A critical feature of the model is to describe how consumers choose products among all competing options, including products offered by our retail partner and its major competitors. In particular, modeling and estimating the substitution across retailers is essential to defining the correct responses to competitors’ price and availability changes.

Our model follows the choice model framework pioneered by Guadagni and Little (1983). This framework has been increasingly applied to studies of consumer behavior and firms’ pricing decision in retailing. For example, Kök and Xu (2011) solve the joint assortment and pricing problem under a nested logit model. Honhon et al. (2010) solves the assortment planning and inventory decisions problem under arbitrary customer preferences and substitution behavior. More recently, Farias et al. (2013) propose a nonparametric approach to predicting revenue from offering a particular assortment of choices. Jagabathula and Rusmevichientong (2016) solve the profit-maximizing assortment and prices using a nonparametric choice model.

What distinguishes our approach from a standard choice model is the incomplete information on choice decisions. In particular, we do not observe choices

文章整体思路

made on competitors' products, a common challenge almost all retailers face. If we did observe these choices, it would be straightforward to apply a standard multinomial logit model or some variation of it to estimate how consumers make choices among all options, where each option is a retailer-product pair. Absent competitors' sales information, the classical approach of incorporating a consumer's purchase of a competitor's product is to model it as part of the outside option (Vulcano et al. 2012, Haensel and Koole 2011). Such an approach, however, does not allow us to explicitly estimate consumers' retailer preferences and the extent to which customers engage in price comparison across retailers. Both components are key to identifying substitution patterns across competing retailers. Though variations in product features such as prices assist identification (Newman et al. 2013), they may provide only weak identification in settings with limited price variations or high price correlation among competitors. Our identification strategy exploits temporary variations in consumer choice sets through our retail partner and competitor stockouts, in addition to competitive price variations.

What further distinguishes our work from this literature is the use of field experiments to demonstrate the strength of the method. We design two experiments: one is designed to correct for price endogeneity in estimating demand models, the other to test the applicability and effectiveness of our methodology in real competitive business settings. We now discuss them in detail.

**1.1.3. Field Experiments to Test Price Elasticities and the Best-Response Pricing Algorithm.** In the retail setting, we are aware of two papers that use field experiments to study the impact of price variations: Gaur and Fisher (2005) and Ferreira et al. (2016). These papers focus on how demand varies with prices for products sold by the retailer, holding competition constant. The key differences between our work and theirs include the presence of competition, stockouts, and the fast-changing online environment, which calls for dynamic responses.

We first conduct a randomized price experiment in the field to obtain unbiased measures of price elasticities, thereby overcoming several problems that arise in the observational data of our retail partner. Specifically, demand can appear inelastic to price due to the lack of price variations, which happens very often when selling millions of products online. It also happened that, even if the price of a product itself varied, it followed competitors' prices closely such that there is little price variation comparatively. In this situation, it is impossible to distinguish how demand responds to changes in one retailer's own price versus changes in competitors' prices.

Moreover, price elasticity estimates can be biased upward when ignoring the fact that retail managers make price decisions based on private demand signals. The effectiveness of commonly used instrumental variables—including cost shifters, average prices in other geographically disparate locations, and lagged prices—tend to be compromised by a combination of factors unique in online retailing. For example, cost shifters do not change nearly as fast as prices for many products, while for other products with almost constant prices, lagged and current prices are highly collinear. Lastly, prices of the same product in other locations are not always available, as some products are not sold in other locations or are off-shelf for the majority of time under study.

It is also important to note that with the presence of competition, price randomization alone will not necessarily guarantee unbiased estimation of elasticities unless competitors' actions are properly accounted for. Ignoring competitors' reactions to our retail partner price changes would bias the estimation because prices can still be correlated with unobserved demand shocks through correlation with competitors' prices. This is why we account for changes in competitors' prices and product availability by using the consumer-choice model.

We then conduct a second field experiment to test the best-response pricing algorithm based on the estimated consumer-choice model. We assign products to two groups to minimize substitution across groups while allowing substitution within each group. We conduct the experiment by alternating the treatment between the two groups in one geographical region, while using two other similar but disparate regions where the retailer operates as comparisons. This design leads to a difference-in-differences-in-differences estimator, which allows us to correct for the potential differences in demand trends between the control and treatment groups using comparison groups subject to similar but independent demand. In the experiment, we exert care in framing and communicating the experiment to the pricing managers such that (i) price managers in other regions are unaware of the ongoing experiment and (ii) the experiment is not framed as a test of an algorithm to replace current practice but rather a decision support tool. The experiment demonstrates an 11% revenue increase while maintaining a margin above a retailer-specified target.

In sum, our methodology has stood the test of a real competitive business environment and demonstrated tangible revenue improvements. Our work helps navigate and evaluate the trade-offs involved in bridging theoretical, empirical, and field work. It thus presents a scalable and replicable methodology to set dynamic prices in online retailing.

总结贡献，可以在竞争商业环境下带来可视化的收入的提高。这种方法可扩展、可复制从而实现线上零售的动态定价

无偏估计很重要

基于顾客选择模型的最优响应定价算法的试验



## 2. Empirical Setting

Our retail partner is a leading Chinese online retailer that originally focused on consumer packaged goods but evolved over time to be a hypermarket. Founded in July 2008, the firm achieved sales of \$1.9 billion in 2013 and was identified by a Deloitte survey in 2011 as the fastest growing technology company in the Asia-Pacific region, with a three-year revenue growth of 19,218%.<sup>1</sup> In our study, we focus on one particular product category: **baby-feeding bottles**.

Pricing decisions are present in every category the retailer offers. We initially focused on one category to allow us to consider all of the different factors affecting the pricing decision. Because experimentation is an important component of our research approach, the category must be one with which the retailer is willing to experiment. For these reasons, the retailer suggested we start with baby-feeding bottles before rolling out to other categories with similar features.

This category presents **attractive features** for our study. It includes a group of relatively homogeneous products that can be characterized by a number of well-defined **product attributes**: country of origin, brand, bottle size, bottle shape and material, nipple size, nipple shape and material, and price point. The fact that feeding bottles have well-defined product characteristics makes it easier to identify competing or substitutable products, which plays a key role in the pricing decision. In addition, although there are innovations and new product launches in the baby-feeding-bottle category, the life cycle of the products is long compared to the time span that the product will be used. It is also the case that during the course of our analysis, there were no new product introductions or innovations. The baby-feeding-bottle category presents a relatively small number of brands and manufacturers that do not engage in exclusivity deals with retailers. This means all competing retailers can carry any product across different brands. Finally, most customers will not engage in repeated purchases in a short period of time (e.g., daily or weekly) since the product will outlast the baby's need. Therefore, it is unlikely consumers will anchor prices based on their purchase histories. Moreover, inter-temporal substitution is not a pressing concern, nor is stockpiling behavior, which may be present for other categories, such as toilet paper and laundry detergent.

Although the characteristics of the baby-feeding-bottle category make it very appealing for our purposes, it is important to note these features are not unique to this category. Many other product categories share similar characteristics to which our methodology also apply, such as small appliances, hardware tools, and kitchenware. Hence, the methodology we introduce can be used to create a broadly applicable pricing tool for a variety of product categories and retail settings.

## 3. Consumer-Choice Model

In this section, we first present the **general framework of our choice model**, which describes how consumers make choices among substitutable products that multiple competing retailers offer. Then, we discuss a **strategy for estimating model parameters in the absence of competitors' sales data**.

### 3.1. Choice Model Framework

Facing a choice set of  $J$  products offered by  $R$  retailers, a **consumer  $i$**  obtains utility  $u_{ijrt}$  from purchasing product  $j$  at retailer  $r$  at price  $p_{jrt}$  on **day  $t$** , or utility  $u_{i0t}$  by not purchasing from one of the  $R$  retailers, where

$$u_{ijrt} = \alpha_j + \alpha_r + \beta_j \log p_{jrt} + \epsilon_{ijrt}, \\ j = 1, 2, \dots, J, r = 1, 2, \dots, R, \\ u_{i0t} = X_{0t} \gamma + \epsilon_{i0t}.$$

A consumer will purchase product  $j$  from retailer  $r$  at price  $p_{jrt}$  if  $u_{ijrt} = \max\{u_{ijrt}, u_{i0t}; j = 1, 2, \dots, J, r = 1, 2, \dots, R\}$ , and will choose the outside option if  $u_{i0t} = \max\{u_{ijrt}, u_{i0t}; j = 1, 2, \dots, J, r = 1, 2, \dots, R\}$ . The intercept  $\alpha_j$  corresponds to the constant utility obtained from purchasing product  $j$  regardless of which retailer the product is purchased from. The intercept  $\alpha_r$  is the additional utility obtained by purchasing a product from retailer  $r$ , which can be understood as a **customer's retailer preference**. For instance, a consumer would assign a higher utility to a retailer who offers more convenient online checkout, a reliable delivery program, or a lenient return policy. The higher the value of  $\alpha_r$ , the larger the premium a customer is willing to pay to buy the product from retailer  $r$ . Note that only the **differences** across these retailers preferences are identifiable. Hence, we **normalize  $\alpha_r = 0$  for our retail partner**. Product-specific price sensitivity is captured by the parameter  $\beta_j$ . We do not explicitly model shipping costs because all retailers offer a generous shipping policy in this context—free shipping for a small minimum spending per order (¥29 to ¥39), thanks to low labor costs in China.<sup>2</sup> However, in a context where shipping costs vary significantly across competitors and orders, one could include shipping cost sensitivity in the model.

The **outside option** in our model includes purchasing from other channels, including both online and brick-and-mortar retailers, and not purchasing at all. We allow the utility of the outside option to vary across days of the week and holiday periods as purchase intention could vary between weekdays and weekends or between holidays and “regular” days (Perdikaki et al. 2012, Lu et al. 2013). These covariates—i.e., day-of-week dummies and holiday dummies—are represented by  $X_{0t}$ . Finally,  $\epsilon_{ijrt}$  represents consumer  $i$ 's utility shock of purchasing product  $j$  at retailer  $r$  on day  $t$ .

选择婴儿奶瓶的原因

婴儿奶瓶这品类的特征

不考虑新产品引入/创新；所有的竞争者可以获得不同品牌的产品；大部分顾客短期内不会重复购买

选择框架：哪个效用大就选择哪个

参数的含义

方法的扩展性

Distribution assumptions and correlation patterns of  $\epsilon_{ijrt}$  will be discussed subsequently.

The majority of prior studies involving consumer-choice models restrict the attention to a model where the parameter  $\beta$  is a constant that does not vary across products. In these models, the estimates of price sensitivity are driven primarily by demand and price variations across products. In our work, however, we conduct an experiment to introduce price variations within each product, thereby allowing us to measure the extent to which price sensitivities vary across products, and meanwhile addressing potential concerns of price endogeneity.

There are several arguments in the literature for why price sensitivity might vary by product. First, there are many examples of price premiums charged for products with higher expected quality. This suggests that either higher expected product quality reduces price sensitivity or that less price-sensitive consumers are drawn to higher-quality products (Erdem et al. 2002). Second, product uncertainty may affect price sensitivity. The direction of the effect can go both ways. When consumers are uncertain about product quality, they may use price as a signal and thus exhibit lower price sensitivity (Gaur and Fisher 2005). On the other hand, if consumers are risk averse, they may derive greater disutility from a given price, thus inducing higher price sensitivity for uncertain products (Tellis and Gaeth 1990). Lastly, availability of alternative choices may lead to higher price sensitivity (Nelson 1974). Hence, products offered at more venues may exhibit higher price sensitivity, and popular products may exhibit higher price elasticity than niche products.

An alternative to letting price elasticity vary by product is to let it vary by customers through a random coefficient model, where the price coefficient  $\tilde{\beta}_i$  is consumer specific and is drawn from a distribution whose parameters are to be estimated. The advantage of this model is that it explicitly incorporates consumer heterogeneity. However, how price elasticity varies across products is dictated by product specific intercepts (see Train 2009, Chap. 6, for details). In contrast, the model with product-specific price elasticity allows for greater degrees of freedom and is more sensitive to demand and price variations associated with each specific product; as we shall see in the estimation results, price elasticities vary significantly across products. Of course, achieving this requires greater price variations within each product to retain the statistical power, which we obtain thanks to our randomized price experiment. A more ambitious model may incorporate both consumer heterogeneity and product specificity at the same time. However, such a model requires estimating at least  $J$  random coefficient distributions (both means and standard deviations), which suffers from over-fitting issues when applied to a relatively sparse experimental data set.

相比之下,具有产品特定价格弹性的模型允许更大的自由度,并且对与每种特定产品相关的需求和价格变化更敏感;正如我们在估算结果中所看到的那样,各个产品的价格弹性差异很大。当然,要实现这一点,需要在每个产品中有更大的价格差异来保留统计能力,这是我们通过随机价格实验得到的结果。更雄心勃勃的模式可能同时兼顾消费者的异质性和产品特性。然而,这样的模型需要至少估计J个随机系数分布(平均值和标准偏差),当应用于相对稀疏的实验数据集,这会出现过度拟合的情况

### 3.2. Extent of Price Comparison

The utility shocks  $\epsilon_{ijrt}$  are not completely independent of each other because the  $R$  options associated with a particular product  $j$  are essentially the same product. Even though the purchasing utility could vary depending on the retailer's platform, the consumption utility associated with these products is the same. Consequently, it is reasonable to assume that a consumer who likes product  $j$  at retailer  $r$  should also like the same product offered by other retailers. In other words, the utility shocks  $\epsilon_{ijrt}$  for a product are likely correlated across retailers. To allow such correlation, we assume the utility shocks  $\epsilon_{it} = (\epsilon_{i0t}, \epsilon_{ijrt}; j = 1, 2, \dots, J, r = 1, 2, \dots, R)$  have a cumulative distribution given by

$$\exp\left(-e^{-\epsilon_{i0t}} - \sum_{j=1}^J \left(\sum_{r=1}^R e^{-\epsilon_{ijrt}/(1-\lambda)}\right)^{1-\lambda}\right).$$

Under such a joint distribution, the marginal distribution of each utility shock  $\epsilon_{ijrt}$  follows a univariate extreme value distribution. In other words, our model establishes a nested structure where each product is a nest. The parameter  $\lambda$  can be intuitively understood as an indicator of correlation for utility shocks for the same product offered by different retailers. As  $\lambda$  increases, the correlation increases.<sup>3</sup> A value of  $\lambda = 0$  indicates no correlation, and the model reduces to a standard multinomial logit model. As the value of  $\lambda$  approaches 1, utility shocks approach perfect correlation, which means that all idiosyncratic utility shocks associated with the same product are identical across retailers. In this case, every consumer will always buy from the retailer that offers the lowest price (assuming for a moment that retailer preferences  $\alpha_r$  are identical). That is, consumers engage in most extensive price comparison when deciding where to buy a product. The larger  $\lambda$ , the more likely prices (or prices normalized by channel quality  $\alpha_r$ ) will drive retailer choice. The smaller  $\lambda$ , the more likely the choice of retailers will be proportional to their market share according to the independence of irrelevant alternatives (IIA) property, which asserts that the ratio of probabilities of choosing two alternatives is independent of the availability or attributes of a third option. For this reason, the larger  $\lambda$ , the more concerned retailers should be about monitoring and following competitors' price movements. Under this proposed model, the probability of purchasing product  $j$  from retailer  $r$  on day  $t$  can be written as follows (Train 2009):

$$\Pr_{jrt} = \left( \exp\left(\frac{\alpha_j + \alpha_r + \beta_j \log p_{jrt}}{1 - \lambda}\right) \cdot \left( \sum_{s=1}^R \exp\left(\frac{\alpha_j + \alpha_s + \beta_j \log p_{jst}}{1 - \lambda}\right) \right)^{-\lambda} \right)$$

在t天从零售商r购买产品j的概率

$$\cdot \left( \exp(X_{0t}\gamma) + \sum_{j=1}^J \left( \sum_{s=1}^R \exp\left(\frac{\alpha_j + \alpha_s + \beta_j \log p_{jrt}}{1-\lambda}\right) \right)^{1-\lambda} \right)^{-1} \cdot (1)$$

### 3.3. Stockouts as a Source of Identification

The model discussed so far could be estimated using a standard nested logit framework (Train 2009) if we were able to observe competitors' sales, which, unfortunately, we are not. However, we do observe the assortment competitors carry as well as their prices and availability. In what follows, we first illustrate the identification challenges that incomplete sales information poses. Then, we show how own and competitors' stockout occasions serve as a useful source to identify retailer preferences and the extent of price comparison. We will use a moment condition framework to illustrate how stockouts provide the information needed for identification; however, the model is estimated using maximum likelihood estimator consistent with what is commonly used for the family of logit models. Moment estimators and maximum-likelihood estimators yield similar results (Berry et al. 1995). When sales levels are low, however, such that market shares in specific instances are not good proxies for theoretical probabilities, the maximum-likelihood estimator is more efficient.

For illustration, we use a simple case where there are only two products, 1 and 2, and two competing retailers, A and B (normalize  $\alpha_A = 0$ ). We also assume for now that all utility shocks are independent, i.e.,  $\lambda = 0$ . For simplicity, we drop subscript  $t$  and remove the covariates matrix  $X_{0t}$  and assume that the mean utility of the outside option equals zero. Thus, the model reduces to a standard multinomial logit model where

$$\begin{aligned} u_{i1A} &= \alpha_1 + \beta_1 \log p_{1A} + \epsilon_{i1A}, \\ u_{i2A} &= \alpha_2 + \beta_2 \log p_{2A} + \epsilon_{i2A}, \\ u_{i1B} &= \alpha_1 + \alpha_B + \beta_1 \log p_{1B} + \epsilon_{i1B}, \\ u_{i2B} &= \alpha_2 + \alpha_B + \beta_2 \log p_{2B} + \epsilon_{i2B}, \\ u_{i0} &= \epsilon_{i0}. \end{aligned}$$

Suppose we only observe retailer A's sales data. Given market size  $M$  and sales  $y_{1A}$ ,  $y_{2A}$ , one can calculate the market share of product 1 at retailer A,  $s_{1A} = y_{1A}/M$ , and the market share of product 2 at retailer A,  $s_{2A} = y_{2A}/M$ . We then let market shares be equal to purchase probabilities obtained through Equation (1). By taking the ratios of these market shares, we obtain the following two moment conditions:

$$\frac{s_{1A}}{s_{2A}} = \frac{\exp(\alpha_1 + \beta_1 \log p_{1A})}{\exp(\alpha_2 + \beta_2 \log p_{2A})}, \quad (2)$$

$$\frac{s_{1A}}{1 - s_{1A} - s_{2A}} = \exp(\alpha_1 + \beta_1 \log p_{1A}) \cdot [1 + \exp(\alpha_1 + \alpha_B + \beta_1 \log p_{1B}) + \exp(\alpha_2 + \alpha_B + \beta_2 \log p_{2B})]^{-1}. \quad (3)$$

By rewriting Equation (2), we have

$$\log\left(\frac{s_{1A}}{s_{2A}}\right) = \alpha_1 - \alpha_2 + \beta_1 \log p_{1A} - \beta_2 \log p_{2A}. \quad (4)$$

From Equation (4), we are able to identify three (sets of) parameters:  $\alpha_1 - \alpha_2, \beta_1, \beta_2$ , where  $\beta_1$  and  $\beta_2$  are identified through variations of two products' prices,  $p_{1A}$  and  $p_{2A}$ . Recall that we are able to identify price coefficients without bias because prices  $p_{1A}, p_{2A}$  are randomly assigned in our experiment and, more importantly, assigned in a way that allows us to identify two separate price sensitivities without encountering multicollinearity among these two price series. Unfortunately, we are unable to separate the intercepts  $\alpha_1, \alpha_2$  from this moment condition.

Since  $\beta_1, \beta_2$  can be identified from Equation (4), we are left with three parameters,  $\alpha_1, \alpha_2, \alpha_B$ , that need to be estimated and we have only two moment conditions given by Equations (2) and (3). We would need at least one more moment condition to fully identify all three parameters. In what follows, we show how stockouts provide this additional moment condition. Suppose that on a different day, product 1 stocks out at competitor B. We then have the following moment condition with regard to the newly observed market shares of product 1 and 2 by retailer A,  $s'_{1A}, s'_{2A}$ :

$$\frac{s'_{1A}}{1 - s'_{1A} - s'_{2A}} = \frac{\exp(\alpha_1 + \beta_1 \log p_{1A})}{1 + \exp(\alpha_2 + \alpha_B + \beta_2 \log p_{2B})}. \quad (5)$$

We now have three moment conditions (2), (3), and (5), and exactly three parameters,  $\alpha_1, \alpha_2, \alpha_B$ , to identify. We then demonstrate how the extent of price comparison  $\lambda$  is identified. We again illustrate the logic with an example employing two products and two retailers, but now utility shocks can be correlated across retailers selling the same product and the correlation is parameterized by  $\lambda$ . Consider again two days: on the first day all products are in stock at both retailers, and on the second day product 1 stocks out at competitor B. Market shares on the first day are denoted as  $s_{1A}, s_{2A}$ , and market shares on the second day are denoted as  $s'_{1A}, s'_{2A}$ . We have

$$\frac{s_{1A}}{s_{2A}} = \frac{\exp\left(\frac{V_{1A}}{1-\lambda}\right) \left[ \exp\left(\frac{V_{1A}}{1-\lambda}\right) + \exp\left(\frac{V_{1B}}{1-\lambda}\right) \right]^{-\lambda}}{\exp\left(\frac{V_{2A}}{1-\lambda}\right) \left[ \exp\left(\frac{V_{2A}}{1-\lambda}\right) + \exp\left(\frac{V_{2B}}{1-\lambda}\right) \right]^{-\lambda}},$$

可以通过试验估计  $\beta_1, \beta_2$

(2, 3, 5) 三式可以估计  $\alpha_1, \alpha_2, \alpha_B$

估计  $\lambda$



$$\begin{aligned} & \frac{s_{1A}}{1 - s_{1A} - s_{2A}} \\ &= \frac{\exp\left(\frac{V_{1A}}{1-\lambda}\right) \left[ \exp\left(\frac{V_{1A}}{1-\lambda}\right) + \exp\left(\frac{V_{1B}}{1-\lambda}\right) \right]^{-\lambda}}{1 + \sum_{j=1,2} \exp\left(\frac{V_{jB}}{1-\lambda}\right) \left[ \exp\left(\frac{V_{jA}}{1-\lambda}\right) + \exp\left(\frac{V_{jB}}{1-\lambda}\right) \right]^{-\lambda}}, \\ & \frac{s'_{1A}}{s'_{2A}} = \frac{\exp(V_{1A})}{\exp\left(\frac{V_{2A}}{1-\lambda}\right) \left[ \exp\left(\frac{V_{2A}}{1-\lambda}\right) + \exp\left(\frac{V_{2B}}{1-\lambda}\right) \right]^{-\lambda}}, \\ & \frac{s'_{1A}}{1 - s'_{1A} - s'_{2A}} = \frac{\exp(V_{1A})}{1 + \exp\left(\frac{V_{2B}}{1-\lambda}\right) \left[ \exp\left(\frac{V_{2A}}{1-\lambda}\right) + \exp\left(\frac{V_{2B}}{1-\lambda}\right) \right]^{-\lambda}}, \end{aligned}$$

where  $V_{jr} = \alpha_j + \alpha_r + \beta_j \log p_{jr}$ ,  $j = 1, 2$ ,  $r = A, B$ . Again, the identification of  $\beta_j$ ,  $j = 1, 2$  comes from the variation in prices  $p_{jA}$ ,  $j = 1, 2$ . These four moment conditions allow us to identify four parameters  $\alpha_1, \alpha_2, \alpha_B, \lambda$ .

It is worthwhile to understand the intuition in addition to the mathematical derivation. When a competitor stocks out, the extent to which the retailer's total sales increase provides information about the competitor's market share, or consumers' preference for the competing retailer  $\alpha_r$ . This is because the higher the competitor's market share, the more overflow we would observe when this competitor stocks out. Additionally, if more of the unsatisfied demand spills over to the same product that stocks out at the competitor, rather than to other products, it suggests that tastes are strongly correlated for the same product offered by different retailers, and that more customers shop across retailers. That is, when a product stocks out at a competitor, customers originally interested in the product will not simply buy a different product but instead will be more likely to buy the same product from other retailers.

To summarize, let  $z_{jrt} = 1$  denote that product  $j$  is in stock at retailer  $r$  on day  $t$ , and  $z_{jrt} = 0$  otherwise. We can rewrite the probability of purchasing product  $j$  at retailer  $r$  on day  $t$ :

$$\begin{aligned} \Pr_{jrt} = & \left[ z_{jrt} \exp\left(\frac{\alpha_j + \alpha_r + \beta_j \log p_{jrt}}{1-\lambda}\right) \cdot \left( \sum_{s=1}^R z_{js} \exp\left(\frac{\alpha_j + \alpha_s + \beta_j \log p_{jrt}}{1-\lambda}\right) \right)^{-\lambda} \right] \\ & \cdot \left[ \exp(X_{0r}) + \sum_{r=1}^R \sum_{j=1}^J z_{jrt} \exp\left(\frac{\alpha_j + \alpha_r + \beta_j \log p_{jrt}}{1-\lambda}\right) \cdot \left( \sum_{s=1}^R z_{js} \exp\left(\frac{\alpha_j + \alpha_s + \beta_j \log p_{jrt}}{1-\lambda}\right) \right)^{-\lambda} \right]^{-1}. \quad (6) \end{aligned}$$

**3.3.1. Comparison with Price Variations as a Source of Identification.** Using logic similar to that above, one could prove that price variations at competitors would also lend additional moment conditions for identification, which is consistent with Newman et al. (2013). However, this identification strategy can sometimes be weak, especially if price variations are low, due either to limited self-price variations or highly correlated prices across competitors. It is because under these conditions, besides availability, price is the only time-variant product feature that is present in this context for identification of retailer preferences and extent of price comparison, and the same variation is also used to identify other parameters—specifically, price elasticities  $\beta_1, \beta_2$ . All of the other product features, such as size, brand, or shape, are time invariant.

It is important to note that although prices and availability are both time-variant product features, they enable identification in different ways. While price variation affects only the utility levels obtained from each option, availability variation alters the composition of the choice set by changing what the set offers. The latter generates sizable changes in the choices available to consumers, and this is especially valuable when market shares are small. On the other hand, price variations usually do not affect choices as significantly as availability variations, and price variations are also used to identify other parameters as well, as previously noted.

Using simulation, we compare parameter estimation accuracy across 36 scenarios with varying levels of price variation, price correlation between the retailer and the market, and stockout probabilities. Details of the simulation are discussed in Section EC.2 of the e-companion. We find that if price variation is low and the retailer's price is highly correlated with that of competition, then ignoring stockouts leads to biased and imprecise parameter estimates.

In a nutshell, stockouts allow us to exploit temporal changes in the consumer choice set. Conlon and Mortimer (2013) also use stockouts as for identification but with complete information on choice decisions. As we have shown in this section, this source of variation is even more critical with partially observed choice decisions.

**3.3.2. Other Considerations for Identification.** First, the identification depends on the exogeneity of stockouts. If we observe correlation between stockouts of a product and the unobserved utility shock of this product, the exogeneity condition would be violated. This would be the case, for example, if a product has been popular for some time and has sold out as a result of this popularity. If utility shocks are serially correlated, then the utility shock of the product is likely to be high during the stockout as well. We test this by examining



the demand patterns of each product for autocorrelation and find no evidence of such.<sup>4</sup>

Second, it is possible that retailers **adjust prices** of other available items when they experience stockouts for some products. If the choice model explicitly accounts for competitors' prices of all major products, like we do, such correlation is not a threat to identification.

Third, in contexts where a nonnegligible number of **customers choose to wait** when experiencing a stockout—for example, when the iPhone was introduced—stockouts can still be used as an identification strategy, but the model needs to explicitly account for the waiting decision. In the current model, if a consumer experiences a stockout at a retailer and decides to wait until the product becomes available, the model will count the consumer as choosing the outside option and treat the individual as a new customer when he or she comes back. In our setting, this is not a major concern because the channels and products are not substantially differentiated, unlike iPhones for example. This means it is unlikely that customers will wait to buy a product at a specific retailer when the identical product is available for a comparable price at a competing retailer.

#### 4. Price Elasticity Field Experiment

As previously noted, a challenge that practitioners and academics encounter in identifying the relationship between prices and demand data is the endogeneity that pricing decisions generate. This **endogeneity** challenge holds when the estimation is done from observational data, where prices are not randomly defined and managers make pricing decisions based on demand information observed at that time. Unfortunately, in most situations, researchers cannot observe this information.

In the past, researchers have used instrumental variables to **correct for this bias**. The **two most commonly used instruments** in settings similar to ours are **Hausman-style instruments** and **lagged prices**. Although theoretically appealing, these approaches can have significant limitations in practice. First, **Hausman-style instruments** use the **average price** of the same product sold in other geographical markets as the instrument for the price in the focal market, as prices in different locations are subject to common firm-level cost shifters but are not correlated with demand in the focal market (Hausman 1996, Nevo 2001). The effectiveness of this approach depends on (i) whether the product is sold at multiple locations and (ii) sufficient correlation between prices and common cost shifters. These assumptions may not be valid in practice. For example, if product assortment varies across locations, Hausman-style instruments may not be present for all products considered. **Another limitation**

can be present for products with prices that change frequently—daily or sometimes multiple times a day. Cost parameters are only weak instruments in this case, since costs do not change nearly as frequently as prices do.

Similarly, to use **lagged prices** as an instrument, we need **sufficient correlation** between lagged and current prices, which is often the case (Tereyağoglu et al. 2017 and Villas-Boas and Winer 1999). However, we also need **sufficient time-series variations in prices** such that lagged and current prices are **not highly collinear**. For those products with prices that hold relatively constant during the period of study, lagged and current prices are almost always the same. Hence, the validity of lagged prices as instruments is challenged. **For products with little historical price variation, it is difficult to estimate product-level price elasticity regardless of what instrument is used.** In section 5, we will illustrate the limitation discussed above by comparing the results from different instrumental variable approaches with our experimental approach.

##### 4.1. Field Experiment Design and Implementation

With the support of our retail partner, we implement a **randomized experiment** to **resolve the endogeneity issue**. The field experiment consists of randomly varying prices for the 15 top-selling products in the baby-feeding-bottle category during a four-week period. These 15 products represent 7% of all the products in the category and capture 54% of the unit sales and 55% of the revenue for the category during an average week before the test was implemented. Table 1 presents summary statistics of these 15 products before the test implementation. Based on the standard deviations and the means, we note that the average within-product price c.v. (coefficient of variance) is 5.9%, while the highest price c.v. is 13.8% and the lowest is 1.2%.

We focus on the 15 top sellers for **two reasons**. First, to estimate the impact of the randomized test, we need to focus on products that **have a daily sales rate large enough to allow us to obtain statistically significant results**. Second, this approach aligns with our ultimate goal of developing a robust methodology to **simultaneously maximize total revenue for a group of products**. Hence, we focus on the products within this category that have the largest impact on the category revenue.

When estimating the price elasticity of each product, we are concerned about **different drivers** that potentially can affect customers' sensitivity to prices: the **absolute price level** and the **relative gap between the retailer's price and competitors' prices**. The assumption is that price elasticity depends on **two main factors**: **absolute price variations** and **price variations with respect to the competition**. To deal with this situation, we implement the field test by randomly assigning the products being tested to two different groups.

lagged prices的局限

实验设计

关注前15种的原因

最终

影响顾客价格敏感性的驱动因素

**Table 1.** Summary Statistics for the 15 Products in the Field Test

Product #	Daily sales (units)				Price (¥)				Market price (¥)
	Average	Max.	Min.	SD	Average	Max.	Min.	SD	Average
1	29.0	63.0	0.0	14.5	35.6	40.4	33.9	1.5	40.5
2	15.2	33.0	0.0	8.6	32.2	35.9	31.9	0.7	35.1
3	2.9	15.0	0.0	2.8	40.6	45.0	37.5	2.6	48.6
4	1.8	6.0	0.0	1.4	44.9	48.0	40.0	1.5	48.7
5	3.5	12.0	0.0	2.2	89.0	109.8	80.0	3.9	91.3
6	4.3	13.0	0.0	3.4	88.3	102.0	84.0	2.9	91.2
7	4.3	18.0	0.0	3.6	76.1	82.0	74.3	0.9	79.2
8	2.2	8.0	0.0	2.1	66.2	84.0	61.2	6.6	81.2
9	9.6	27.0	0.0	7.7	91.7	99.0	82.4	6.5	92.8
10	3.7	9.0	0.0	2.5	14.4	16.7	13.9	0.6	20.9
11	17.8	36.0	0.0	7.2	84.9	108.0	78.9	6.3	100.0
12	22.5	65.0	0.0	18.4	86.0	175.0	83.2	11.9	83.9
13	18.2	58.0	0.0	15.4	85.5	109.0	84.0	5.4	84.0
14	3.9	18.0	0.0	4.3	121.5	134.0	107.0	10.2	121.0
15	2.3	27.0	0.0	5.5	130.7	145.7	118.5	11.7	111.0

Note. These summary statistics correspond to the two-month period right before the field test was implemented.

The first group, **group A**, randomly changes prices with respect to a baseline price defined by the product's historical price at the retailer, while the second group, **group B**, randomly changes prices with respect to the competitors' prices for that product. We focus on the top four competitors that, together with our partner, capture more than 80% of the online B2C retail market in China. Once a product is randomly assigned to one of the two groups (variation with respect to the historical baseline or to the competition price), we randomly assign the treatment level for each of the products on each day, depending on the group to which the product belongs.

We define **five different treatment levels**: high (+10%), medium-high (+5%), medium (no variation), medium-low (−5%), and low (−10%). High and low are defined for each product based on its assigned group. For example, if a product belongs to group A and is randomly assigned “medium-low” treatment, the product price is set 5% below the average price for that product over the eight weeks immediately prior to the beginning of the test. If a product belongs to group B and is randomly assigned “high” treatment, the product's price is set 10% above the lowest competitor's price for that product during that day.

Before running the experiment, we analyzed historical data and simulated outcomes we may possibly obtain from a field test based on assumed parameters. A major consideration in this stage, both internally within the research team and when interacting with the retail partner, was the duration of the test. We estimate the minimum price variation and test duration needed for the field test to be effective. With these numbers, we ensure management that we were not asking for unnecessary efforts on their end. Another issue considered at this stage was the need for the

company to keep all other factors constant during the testing period—for example, no deviations in promotion strategies. The combination of these factors led to a joint decision of 30-day duration and the price variation range that we considered in the test.

We randomly assign treatment using a **fractional factorial design** (Mead et al. 2012), taking into account the five different treatment levels. The treatment assigned to a product remained constant during a three-day period. Note that for products in group B, the treatment is constant, but the actual price could vary during the three-day period as competitors may change their prices. We choose to maintain the treatment level during a three-day period to avoid the risk of alerting customers that a test is in place and to minimize the chance of customers' speculative behaviors.

Each day at 9:00 A.M. local time during the four weeks of the test, we monitored competitor prices and adjusted our own prices according to the randomized schedule. Additionally, we monitored competitive responses by scraping prices and product availability information from competitors' websites 12 hours after our retail partner changed prices each day. In particular, product availability information was obtained by visiting product detail pages at competitors' websites. When a product was out of stock, consumers were still able to find it from a retailer's website but saw a label of “Out of Stock” or “Add to Wishlist” on the search results and product details pages.

#### 4.2. Field Test Analysis and Results

The test lasted for a total of 30 consecutive days and the implementation occurred according to plan. As an illustration, Figure EC.1 in the e-companion presents a comparison between units sold and retail prices for two products included in our sample for the period

两个分组

五个变化水平

部分因子设计

避免顾客投机

连续的

before and after the test was implemented. It is interesting to observe that, while it is very hard to see any pattern between price and demand before the test was implemented (e.g., the top figure for product 1), after random treatments are introduced, the price elasticity becomes evident by visual inspection (graphs at the bottom of the figure). In addition, the field test allows us to introduce price variation on a number of products that historically present a very stable price pattern (e.g., product 2 in the figure). It is evident that if prices do not vary, it is almost impossible to estimate their impact on demand.

Of course, beyond the visual examination of the test results, we want to take full advantage of the randomized experiment. In the next section, we estimate the customer-choice model introduced previously. We then calculate optimal prices that account for, among other controls, the unbiased price elasticity estimates as well as the competition's prices and product availability.

## 5. Estimation Results

We estimate the demand model using the maximum-likelihood estimator which maximizes the overall probability of observing a certain number of purchases for all SKUs on all days during the 30-day test period. The exact formulation is expressed as follows:

$$\max_{\alpha, \beta, \gamma, \lambda} LL(\alpha, \beta, \gamma, \lambda) = \sum_{t=1}^{30} \left[ \sum_{j=1}^J D_{jrt} \log \Pr_{jrt} + \left( M_t - \sum_{j=1}^J D_{jrt} \right) \log \left( 1 - \sum_{j=1}^J \Pr_{jrt} \right) \right].$$

The purchase probability  $\Pr_{jrt}$  is defined in Equation (6), where  $r = 1$  represents our partner retailer;  $D_{jt}$  is the observed sales for product  $j$  on day  $t$  at our partner retailer; and  $M_t$  is the market size on day  $t$ .<sup>5</sup>

Table 2 compares estimation results obtained from different types of data, observational versus experimental, and different types of models, with and without instruments, with and without competition and price comparison. We apply a standard multinomial logit model, to the historical observational data, without instruments (column (1)) and with three instrumental variables (IV; columns (2)–(4)) commonly used in the literature: (i) SKU-level costs (including wholesale costs and operating costs); (ii) Hausman-style instruments (i.e., average prices in other geographically disparate locations (Hausman 1996)); and (iii) lagged prices (Villas-Boas and Winer 1999, Tereyağoglu et al. 2017). Note that these models only account for prices and availability of substitutable products at our partner retailer and do not yet include competitive information. We then apply the same model to the experimental data without using any instrument (column (5)), while in column (6) we include competitive prices and

product availability in the model, and in column (7) we further include consumer extent of price comparison across retailers.

### 5.1. Observational Data vs. Randomized Price Experiment

As noted before, managers may adjust prices based on demand signals unobservable to us, and this may cause an upward bias in price elasticity estimates. This is evident in Table 2; the elasticities of five of the 15 products turn out to be positive when simply applying a multinomial logit model to historical data without accounting for potential price endogeneity (column (1)). Even though positive price elasticity can be found for conspicuous consumption, it would be hard to imagine that this is the case for the product category under study (baby-feeding bottles).

The power of this instrumental variable approach critically depends on (i) the availability of a valid instrument and (ii) whether the instrument itself captures sufficient variations from the original variable to avoid insignificant estimates arising from weak instruments. In our context, we find that the performance of some commonly used price instruments suffer these concerns.

First, when SKU-level costs are used as an instrument in our setting (column (2)), six of the 15 elasticities turned out to be positive, and all estimates are much noisier (as shown by large standard errors) than those in our base analysis without instruments (column (1)). This is because costs do not change nearly as fast as prices in this context (i.e., daily for some products), and for this reason, it turns out to be a weak instrument. Second, even though the Hausman-style instrument, average prices in other geographical locations (column (3)), has the best performance of all three instruments when available, it is subject to limitations as well. The instrument does not exist for some SKUs, because these SKUs are not offered at other locations or are off shelf for most of the time during the study period (available for less than a week during two months). Lastly, results from using lagged prices as instruments are almost identical to the results from using original prices. This is driven by the high correlation between prices and lagged prices. Recall that for some products, prices hardly varied during the period under study.

We now turn to the results based on the randomized price experiments presented in column (5). We note that 13 of the 15 price elasticities are negative, and 10 of the 15 products present a lower price elasticity compared to the one obtained when using historical data. It is worthwhile noting that this is achieved with only half the sample size (30 days versus 67 days). This analysis suggests that with the randomized experiment, we are able to correct, to a certain extent, the positive

价格内生性风险——正的价格弹性

使用历史数据的结果

使用随机价格实验的数据结果与好处——纠正偏差 (by 价格内生性)



**Table 2.** Estimates of Price Elasticities for Feeding Bottle SKUs

	Historical data				Randomized price experiment		
	(1) No comp. info	(2) No comp. info IV cost	(3) No comp. info IV Hausman	(4) No comp. info IV lagged price	(5) No comp. info	(6) W/comp. info ind. demand shocks	(7) W/comp. info corr. demand shocks
Product-specific price elasticity ( $\beta_i$ )							
Product 1	0.974 (0.624)	-4.969 (5.259)	0.060 (0.704)	0.755 (0.664)	-5.496*** (0.939)	-5.438*** (0.234)	-1.675*** (0.157)
Product 2	-2.451*** (0.519)	-10.453 (6.639)	2.328 (2.157)	-2.288*** (0.549)	-2.390** (1.129)	-2.267** (1.004)	-0.367*** (0.014)
Product 3	-5.436*** (1.043)	7.190 (13.222)	-5.200** (2.561)	-8.066*** (1.399)	-7.109*** (1.715)	-4.572*** (0.636)	-6.773*** (0.643)
Product 4	-4.398** (1.8216)	-5.068 (5.0866)	-1.511 (6.7204)	-2.648 (2.1419)	2.280 (2.0565)	-0.006 (0.025)	-0.004 (0.001)
Product 5	-1.848** (0.906)	-0.525 (5.678)	-3.193 (2.381)	-1.161 (1.228)	0.096 (2.468)	-0.004*** (0.003)	-0.953*** (0.103)
Product 6	-1.410 (1.213)	-10.999 (10.129)	-1.779 (2.797)	-0.796 (1.261)	-4.028* (2.332)	-2.367*** (0.523)	-1.054 (0.727)
Product 7	-0.214 (0.649)	2.076 (9.952)	-3.764 (2.666)	-0.252 (0.714)	-1.534 (1.415)	-2.399*** (0.472)	-0.540 (0.340)
Product 8	3.709 (2.422)	10.121*** (2.467)	N/A N/A	5.768** (2.402)	-4.222*** (1.406)	-2.964*** (0.336)	-1.164** (0.535)
Product 9	-3.713*** (0.699)	-11.545*** (2.020)	-5.500*** (1.269)	-3.587*** (0.712)	-4.450*** (0.872)	-4.296*** (0.293)	-1.118*** (0.159)
Product 10	2.244 (1.506)	-0.684 (10.976)	-0.363 (2.475)	3.001* (1.640)	-4.400 (3.381)	-4.845*** (0.551)	-4.149*** (1.343)
Product 11	-8.225*** (0.810)	-9.944*** (1.090)	-7.994*** (0.817)	-8.279*** (0.846)	-4.324*** (0.959)	-4.058*** (1.872)	-0.504*** (0.093)
Product 12	1.487** (0.654)	2.151 (5.002)	-4.919*** (1.606)	-0.795 (0.611)	-17.149*** (5.917)	-3.049*** (0.077)	-2.187*** (0.045)
Product 13	0.812 (1.252)	-21.701*** (7.294)	-26.739*** (5.726)	1.505 (1.435)	-26.061*** (3.075)	-4.515*** (0.284)	-11.281*** (0.518)
Product 14	-3.632*** (0.546)	7.548 (13.609)	N/A N/A	-3.438*** (0.519)	-2.110* (1.151)	-2.525*** (0.209)	-0.922*** (0.167)
Product 15	-2.914*** (0.759)	10.459 (16.980)	N/A N/A	-2.821*** (0.703)	-3.460*** (0.709)	-3.848*** (0.197)	-1.142*** (0.151)

biases caused by price endogeneity.<sup>6</sup> In addition, the results not only show improvement over the historical data without instruments, but also over historical data with commonly used instruments.

## 5.2. With vs. Without Competitor Information

We next focus our attention on how our proposed method, which accounts for competitor prices and availability (column (6)), performs relative to the standard multinomial logit model (columns (1)–(4)). It is important to emphasize that in our setting, price randomization per se is no guarantee for exogeneity. This is because if competitors respond by adjusting their prices based on the randomized prices, not accounting for competitors' actions would introduce correlation between our prices and the unobserved demand shocks through correlation with competitor prices. If competitor prices and our prices are positively correlated, then there will be a positive bias in the estimation

of price sensitivities,  $\beta_i$ , unless we account for competitors' prices appropriately. To illustrate this, suppose competitors follow our prices closely and decrease their prices as we decrease ours; as a result, we will receive a lower demand shock. This generates a positive correlation between our price and our demand and hence biases price elasticity estimates upward.

To explore this issue, we study whether it is the case that, during the period of our experiment, the competitors react to our randomized changes by changing prices accordingly. We first note the percentage of time that daily prices changed from the previous day: 38.6% (our retail partner), 1.9% (competitor 1), 16.3% (competitor 2), 19.5% (competitor 3), and 11.1% (competitor 4). Table EC.1 in the e-companion further shows the frequencies of competitor price responses within 12 and 24 hours after our price change. We see that competitors did not seem to follow closely our random

外生性是否是事实：对手根据我们的随机价格变动改变自己的价格

Table 2. (Continued)

	Historical data				Randomized price experiment		
	(1) No comp. info	(2) No comp. info IV cost	(3) No comp. info IV Hausman	(4) No comp. info IV lagged price	(5) No comp. info	(6) W/comp. info ind. demand shocks	(7) W/comp. info corr. demand shocks
Retailer preferences							
Retail partner						Baseline	Baseline
Competitor 1						0.260 (3.163)	0.217** (0.096)
Competitor 2						0.366 (4.225)	0.017 (0.119)
Competitor 3						−3.750 (10.683)	−1.836*** (0.045)
Competitor 4						−3.887*** (1.457)	−2.464*** (0.010)
Extent of price comparison ( $\lambda$ )							0.7911*** (0.0203)
Product intercepts	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Day-of-week and holiday	Yes	Yes	Yes	Yes	Yes	Yes	Yes
# days	67	67	67	67	30	30	30
# purchases	7,690	7,690	7,298	7,580	3,742	3,742	3,742
Pseudo $R^2$	0.6696	0.6671	0.6749	0.6668	0.5834	0.7107	0.7164
MAD (daily)	0.439	0.436	0.496	0.449	0.451	0.387	0.377
Log likelihood	−35,973.844	−36,650.168	−29,056.259	−36,681.365	−17,281.855	−17,053.73	−17,036.01

Notes. comp.=competitor; ind.=independent; corr.=correlated. Columns (1)–(4) display estimates using historical data without and with instruments. Columns (5)–(7) display estimates from the 30-day randomized price experiment. Column (2) uses cost as the instrument to price. Column (3) uses the Hausman-style instrument, average prices of the same product sold at other geographical locations by the same company. Because certain products are not sold at either of the other locations or off shelf for most of the time during the study period (available for less than one week during two months), price elasticity for these products are not available. Column (4) uses one-day lagged price as instrument. In columns (1)–(6), elasticity $_{jA} = \beta_j(1 - s_{jA})$ , and  $\approx \beta_j$  if  $s_{jA}$  is small. In column (7), elasticity $_{jA} = (\beta_j/(1 - \lambda))(1 - s_{jA}) - (\lambda/(1 - \lambda))\beta_j(s_{jA} + ns_{jA})$ , where  $ns_{jA}$  denotes the share of product  $j$  offered by retailer  $A$  within the product nest—i.e., the set of product  $j$ 's available at all retailers.

\*\*\* $p < 0.01$ ; \*\* $p < 0.05$ ; \* $p < 0.1$ .

price changes. The most responsive competitor, competitor 2, followed only 11 of our 91 price increases and 9 of our 75 price decreases. This pattern is consistent with that in the pretest period.

This result also indicates that during the one-month test period, the experiment successfully introduced random variations not only in our own-price levels, but also in the *relative* price levels to competitors. Using the full choice model, we observe **a significant improvement in the model fit**: the  $R^2$  increased from 0.5848 (column (5)) to 0.7107 (column (6)), a 21.5% increase.<sup>7</sup>

### 5.3. Extent of Price Comparison

Finally, we compare the last two columns in Table 2. Column (7) introduces an additional parameter,  $\lambda$ , that captures **the extent to which consumers conduct price comparisons**. Recall that  $\lambda = 0$  means utility shocks for the same products across retailers are *independent*, in which case columns (6) and (7) should be identical. The larger the value of  $\lambda$ , the higher the correlation of utility shocks for the same product across retailers and the more intense the price comparison. The estimate of  $\lambda$  equals 0.7911 and is significant with

99% confidence, which suggests a nonnegligible correlation in utility shocks.<sup>8</sup> In other words, consumers engage in extensive price comparisons across competitors. The high intensity of price comparison indicates that retailers need to follow competitors' price adjustments closely to stay competitive in the market. The impact of either overpricing or underpricing can be significant. In fact, despite the small difference in overall model fit, columns (6) and (7) will suggest very different responses when competitors change prices, as we later discuss.

We now turn our attention to the estimate of competitor significance, i.e.,  $\alpha_r$ . Our retail partner is the baseline. Competitor 1 admits the highest preference, followed by competitor 2, our retail partner, and competitors 3 and 4, respectively. This is consistent with management knowledge of the market.<sup>9</sup> Except for competitor 2, all estimates are statistically significant (as compared to the baseline). All of these retailers sell directly to customers. We do not account for prices charged by small individual sellers who sell through online platforms, because no single seller sells enough to be considered a major competitor, and these small

6、7除了R2接近，响应结果有很大不同

顾客偏好


λ取值不同代表的含义

indeed close to our partner retailer. Nevertheless, it would be better if competitors experienced more stock-outs in the month during our experiment.

Figure 1 shows the goodness of fit graphically. It plots the predicted daily sales (light lines) against observed daily sales (dark lines) based on our estimation results in column (7) of Table 2. The average daily mean absolute deviation (MAD) is 0.377. Note that there is a negative correlation of  $-0.697$  ( $\text{sig} = 0.0041$ ) between model goodness of fit by product, as measured by MAD, with average daily sales. In other words, the model better predicts demand for fast-moving products than for slow-moving products.

Based on estimated model parameters, we calculate sample own- and cross-price elasticities as shown in Table 3. Note that own-price elasticity differs by product due to differences in price sensitivities measured by  $\beta_j$ . Also note that our demand is most sensitive to prices of competitors 1 and 2, but not to prices of competitors 3 and 4. This is because competitors 3 and 4 are relatively small players in this category, as revealed by the estimated retailer preference  $\alpha_r$ . However, retailer preference  $\alpha_r$  alone does not completely explain the difference in cross-price elasticity across retailers. Note that sometimes our demand is more sensitive to prices of competitor 2 than those of

拟合优度





**Table 3.** Own- and Cross-Price Elasticities

Product #	Own	Competitor 1	Competitor 2	Competitor 3	Competitor 4
1	-5.5378	-1.2071	-2.8775	-0.0055	-0.0001
2	-1.7681	-0.7598	-0.6386	-0.0012	0.0000
3	-5.4942	-0.0018	-0.0095	-0.0120	-0.0001
4	-0.0046	-0.0093	-0.0069	0.0000	0.0000
5	-1.5826	-0.4744	-0.7552	-0.0013	0.0000
6	-2.5504	-0.7253	-1.2292	-0.0020	-0.0001
7	-0.9213	-0.4088	-0.3209	-0.0006	0.0000
8	-3.6766	-1.8118	-1.0456	-0.0068	0.0000
9	-3.4141	-0.8532	-1.7617	-0.0023	-0.0001
10	-1.8954	-0.0883	-0.0164	-0.0069	0.0000
11	-2.4377	-0.9699	-0.9174	-0.0023	-0.0001
12	-8.2826	-1.5770	-4.9116	-0.0064	0.0000
13	-23.6245	-0.0152	-14.2382	-0.0138	-0.0022
14	-3.3974	-1.6779	-0.9875	-0.0051	-0.0001
15	-4.1404	-1.3791	-1.6345	-0.0094	-0.0001

competitor 1, even though competitor 2 is less preferred than competitor 1 ( $\alpha_2 < \alpha_1$ ). This is because for some products, competitor 2 actually charges a lower price than competitor 1 and thus increases its competitiveness in the market. In sum, which competitors to follow closely depends not only on the average retailer preference, but also on competitors' actual price levels for the SKU being considered.

Comparing results obtained from typical reduced-form linear demand models and our results, we note that, thanks to the randomized price experiments, we are able to obtain accurate measures of price elasticities with simple linear demand models. However, even though we may be able to randomize our own prices, we cannot randomize competitors' prices. Therefore, estimating the effect of competitors' prices on our demand nonparametrically with a reduced-form regression model can be challenging, especially with a short sample period (i.e., 30 days). This is why modeling of consumers' choices of competing retailers becomes appealing. See Section EC.3 in the e-companion for a detailed discussion.

We note as well that estimated price elasticities indeed vary across products. Given that we have only 15 SKUs, a small sample size, regressing estimated price elasticities on product features will not yield statistically significant results. However, we make the following interesting observation on the difference of price elasticities of domestic and imported brands. Among the 15 SKUs, we have five different brands: Pigeon (four SKUs, domestic), Love (one SKU, domestic), NUK (six SKUs, origin Germany), Avent (three SKUs, origin United Kingdom), and Brown (one SKU, origin United States). We find that the three SKUs with the most elastic demand are all domestic products (all Pigeon, price levels of ¥84, ¥84, and ¥36). We also note that the three SKUs with the least elastic demand are all imported brands (two NUK and one

Avent, with price levels of ¥90, ¥76, and ¥44). That is, consumers are less elastic to imported brands than to domestic brands, given similar prices. There can be several possible explanations. In China, imported brands are typically perceived as being of high quality (Ozer et al. 2014), particularly for baby-and-mom products, given the recent scandals in China involving tainted infant milk (Gough 2014). The observation of lower elasticity on prices of imported brands can therefore be explained by higher expected product quality reducing price sensitivity, or less-price-sensitive consumers drawn to higher-quality products such as imported brands (Erdem et al. 2002). It may also be explained by the fact that more channels sell domestic brands than imported brands, as imported brands are typically distributed only through online channels.<sup>10</sup> Products offered at more venues exhibit higher price elasticity (Nelson 1974).

In the next section, we use the estimates obtained from the choice model presented in column (7) of Table 2 to determine the best-response prices for each product.

## 6. Best-Response Pricing

Based on the estimates obtained from the choice model and taking into account competitor prices and product availability, we optimize prices for our retail partner to maximize total revenue for the category subject to several constraints imposed by the retailer. The constraints include a lower bound on average category margin, lower and upper bounds on individual product margins, and manufacturer price restrictions. Changes in recommended prices come from four different sources: (i) changes in costs, (ii) changes in the retailer's own product availability, (iii) changes in competitor prices, and (iv) changes in competitor product availability. On

进口产品价格弹性低

a particular day  $t$ , we solve the following constrained nonlinear optimization problem:

$$\begin{aligned} \max_{p_{1t}, p_{2t}, \dots, p_{jt}} \quad & \sum_{j=1}^J p_{jt} s_{jt}(p_{jt}, z_{jt}; p_{-jt}, z_{-jt}; \mathbf{p}_{ct}, \mathbf{z}_{ct}; \alpha, \beta, \gamma, \lambda) \\ \text{s.t.} \quad & \frac{\sum_{j=1}^J (p_{jt} - c_{jt}) s_{jt}(p_{jt}, z_{jt}; p_{-jt}, z_{-jt}; \mathbf{p}_{ct}, \mathbf{z}_{ct}; \alpha, \beta, \gamma, \lambda)}{\sum_{j=1}^J p_{jt} s_{jt}(p_{jt}, z_{jt}; p_{-jt}, z_{-jt}; \mathbf{p}_{ct}, \mathbf{z}_{ct}; \alpha, \beta, \gamma, \lambda)} \\ & \geq \text{margin target}, \\ & \text{margin } LB_j \leq (p_{jt} - c_{jt})/p_{jt} \leq \text{margin } UB_j, \quad \forall j, \\ & LB_j \leq p_{jt} \leq UB_j, \quad \forall j, \end{aligned}$$

where  $s_{jt}(p_{jt}, z_{jt}; p_{-jt}, z_{-jt}; \mathbf{p}_{ct}, \mathbf{z}_{ct}; \alpha, \beta, \gamma, \lambda)$  denote the market share of product  $j$  at our retail partner, as a function of its own price and availability  $p_{jt}, z_{jt}$ , prices and availability of the other substitutable products at our retail partner  $p_{-jt}, z_{-jt}$ , prices and availability of all products at competitors  $\mathbf{p}_{ct}, \mathbf{z}_{ct}$ , and all of the estimated parameters  $\alpha, \beta, \gamma, \lambda$ . The market share can be calculated using Equation (6).

Hanson and Martin (1996) show that logit profit functions are not concave in prices. One can easily show that our constrained optimization problem is also nonconcave in prices. Therefore, to find the optimal solution, we seed 10 different initial values using the MATLAB optimization procedure for constrained optimization and choose the one that yields the highest revenue. The optimization procedure used to choose the optimal prices for a given day takes approximately

30 seconds on a 16GB RAM computer. We correlated recommended SKU margins with SKU-specific price elasticities and found the correlation to be 52%–57% on different days, suggesting that our algorithm indeed recommends a higher price (relative to cost) when the product's demand is less price elastic (note that less elastic means higher values), which is intuitive.

We then demonstrate the importance of considering correlated demand shocks within a product nest. Table 4 shows how two models, with independent and correlated demand shocks across retailers, lead to different responses when competitors change prices. Note that in this table, inputs to the optimization model (costs, margins, and bounds) are altered for confidentiality reasons. However, the structure of the optimization itself and demand parameters previously estimated ( $\alpha, \beta, \gamma, \lambda$ ) are intact. With independent demand shocks, prices are driven primarily by costs and price elasticities that are product specific. When competitors change prices, lost demand will be allocated to all available alternatives—this means all products at all retailers—proportionally based on their current market shares. Since there is a large number of options, the spillover effect to the same product at a different retailer is very small, and for this reason, the suggested price change is very small (e.g., when competitors' prices change by ¥10, we only change the price of the same product by ¥1 as shown in column (4) of Table 4).

Substitution patterns are key to responsive pricing. At one extreme, if customers do not substitute

**Table 4.** Best-Response Prices Under Different Models

Product #	Competitor 1 price		Suggested own-price independent demand shocks		Suggested own-price correlated demand shocks	
	(1) Day 1	(2) Day 2	(3) Day 1	(4) Day 2	(5) Day 1	(6) Day 2
1	42.8	42.8	38.1	37.4	39.1	39.2
2	36.8	36.8	38.6	38.4	37.5	36.8
3	48.8	48.8	46.8	47.8	45.0	45.0
4	48.8	48.8	49.8	48.3	48.3	48.5
5	108.0	108.0	87.7	87.8	87.5	86.3
6	108.0	108.0	87.4	87.3	89.3	87.9
7	79.0	79.0	81.7	81.5	82.6	82.6
8	85.0	85.0	72.3	72.3	74.4	74.2
9	108.0	108.0	84.4	84.4	87.9	86.6
10	18.0	18.0	17.0	16.7	18.9	18.9
11	108.0	108.0	100.6	100.6	101.5	101.5
12	98.0	98.0	79.1	80.3	81.0	81.4
13	98.0	98.0	80.9	81.3	83.2	83.4
14	105.0	95.0 (↓10.0)	98.8	97.8 (↓1.0)	103.1	97.5 (↓5.6)
15	109.0	119.0 (↑10.0)	104.7	105.9 (↑1.2)	99.3	103.2 (↑3.9)

*Notes.* In this table, we illustrate our reactions in a real scenario where competitor 1 reduces the price for product 14 by ¥10, while raising the price for product 15 by ¥10 from day 1 to day 2. All other competitors' prices remain unchanged. We demonstrate that using the demand model with independent demand shocks versus the demand model with correlated demand shocks yields very different price reactions. With independent demand shocks, the algorithm suggests only a small change in our retail partner's price in response: a decrease of only ¥1.0 for product 14 and an increase of ¥1.2 for product 15. With correlated demand shocks, however, the algorithm suggests more sizable changes in prices: a decrease of ¥5.6 for product 14 and an increase of ¥3.9 for product 15.

across retailers, there is no need to follow competitors' prices. At the other extreme, if customers always compare prices, one should almost always follow competitors' price changes. From our model's estimation, consumers exhibit a strong price comparison behavior ( $\lambda = 0.7911$ ). This explains why under such a model, prices are more responsive to competitor price changes, as shown in column (6) of Table 4; our suggested prices change by a similar order of magnitude when competitors change their prices (e.g., when competitor prices change by ¥10, the estimated prices change by ¥4 to ¥6). We do not match fully the price change of ¥10 because we already price lower than competitors prior to this change.<sup>11</sup>

In sum, a model that fails to capture the extent to which consumers compare prices across retailers will lead to an insufficient response to competitor price changes. This is true even when, as in our setting, we are able to capture unbiased price elasticities with a randomized price experiment.

To gauge the potential revenue impact of our proposed best-response pricing strategy, we conducted an out-of-sample test on the historical data—the same period that we used to compare the estimation results in Table 2. We find that, had the retailer used this best-response pricing algorithm to adjust prices in response to price changes in the market instead, they would have obtained 7% higher category revenue while holding gross margin on a par with the existing practice. The revenue improvement is estimated as follows. First, we obtain the optimal prices by solving the best-response price optimization for the historical period. Second, we obtain daily total unique web traffic to the feeding-bottle category and use it as a proxy for the market size of each day. Note that this measure of market size accounts for both those customers who purchased and those who visited but did not purchase. Lastly, for each customer in the market, we predict her probabilities of purchasing each product on each day given the new prices calculated in the first step. We then obtain expected demand for each product on each day multiplying these probabilities with the market size obtained in the second step.

Finally, we would like to note that when determining the best prices, our methodology factors in only the current prices charged by competitors, and not competitors' potential responses to our price changes (Li et al. 2017).

## 7. Testing Best-Response Pricing with a Controlled Experiment

We test the performance of our best-response pricing algorithm with a field experiment at our partner retailer, in which prices are changed daily according to the outcome of our model.

### 7.1. Experiment Design and Implementation

Note that the objective of our pricing algorithm is to maximize revenue for the category. A specific product's price changes will lead to revenue changes not only for that product, but also for other products with similar features due to substitution. For this reason, a valid experiment design requires minimal substitution between treatment and control groups; otherwise, the control group will be contaminated due to the spillover effect. Instead of matching products on their main attributes, we identify one existing attribute that allows a clean separation between market segments: bottles designed for ranges of babies' ages. Each feeding bottle is designed for a specific age group because babies of different ages require different nipple sizes, nipple shapes, and bottle volume. There is hardly any substitution between feeding bottles that are designed for a particular age group. Within our 15 feeding bottles, we identify nine bottles that are designed specifically for babies ranging from birth to six months old, and the remaining six bottles are designed for babies seven months and older.

Although we believe it is highly unlikely that the pricing algorithm could work for one group of bottles and not for the other, to ensure that the difference between the two age groups does not drive the result of our pricing algorithm, we rotate the implementation of the treatment between these two groups. Table 5 shows the design of the experiment, which lasts for a total of five weeks. Note that in the last week, we let the experiment return to a state where neither of the two groups receives treatment, to further confirm that we are not capturing an overall time-trend effect.

Typically, one can use the difference-in-differences (DiD) estimator to compute the treatment effect in such settings. The critical identifying assumption underlying the DiD estimator is the existence of a parallel trend. That is, the two groups would otherwise follow the same trend in the absence of the treatment. However, in our context particularly, because the control and treatment groups are not similar products matched according to product attributes but rather dissimilar products to minimize across group spillover, it is more likely that the parallel trend assumption may not hold. It is

**Table 5.** Experiment Design: Test Best-Response Pricing

	Group 1 (Baby age: zero to six months)	Group 2 (Baby age: seven months and above)
Week 0	Control	Control
Week 1	Treatment	Control
Week 2	Control	Treatment
Week 3	Treatment	Treatment
Week 4	Control	Control

*Notes.* Control: current pricing practice. Treatment: best-response pricing algorithm.



possible that in a short period of time, the demand for the two groups of products actually follows different trends due to such reasons as seasonality of births. Indeed, we find evidence from data prior to the experiment that products in group 1 exhibit a slight upward trend in demand, while products in group 2 exhibit a slight downward trend in demand, and these trends exist across all regions in which the company operates. When the parallel trend assumption is violated, the DiD estimator will be biased (Wooldridge 2010).

To seek a solution to this problem, we note that our retail partner operates in multiple regions of the country. These regions are geographically separated, 800–1,000 miles apart. Each region has its own warehouse, logistics team, and management personnel. The population within each region is very dense; thus, the demand is sufficiently large to allow for dedicated shipping in which a warehouse ships only to customers within its own region. This means each region can be viewed as a separate market since there is no demand or supply substitution across these markets. The retailer's website requires customers to choose their locations before being able to browse any content. Once a customer places an order, the individual's shipping address will be verified again to avoid any cross-selling or overlap between transactions across regions. This is a common practice among all major retailers in the country.

This modus operandi gives us an advantage in our estimation approach since we are able to introduce an additional comparison group. We implement the pricing algorithm in only one of the regions, region A, leaving the other regions as controls. We match region A with the other regions in the country by looking at online shopping traffic, sales volume, demographics (population density and income distribution), and, in particular, margin rates, one of the key constraints of our optimization procedure. After the process, we identify two regions, B and C, that match region A.

To ensure the validity of regions B and C as controls, we confirm that none of their pricing managers or product managers are aware of the concurrent experiment conducted in region A. In region A, where the experiment is implemented, it is not feasible to keep managers entirely unaware of the experiment because we need their cooperation to be able to adjust prices. However, we exercise extra care not to alter pricing managers' regular decisions so that controls in Table 5 are valid. To do this, we communicate to the managers that they should make pricing decisions as they normally do. Specifically, instead of framing the experiment as a test of a potentially superior pricing algorithm (i.e., a pilot or an implementation), we communicate to the team that the test is a new randomized pricing experiment. Furthermore, when a group

of products receive treatment, we ask a designated person to update the product prices instead of sending the recommended prices to the pricing team, in an effort to avoid noncompliance or biasing their decisions.

## 7.2. Triple-Differences Estimator

The experiment treatment we describe allows us to adopt a triple-differences estimator to measure the impact of the proposed pricing methodology. The triple differences come from the comparisons of periods when the treatment is applied and when it is not; the two age groups; and regions A, B, and C. Note that the key distinction of this estimator from the more commonly used DiD estimator is the presence of a third comparison group (i.e., across regions).

If the null hypothesis is that without treatment the ratio of group 1 revenue (or group 2 revenue) at region A over group 1 revenue (or group 2 revenue) at region B is constant, then we could use the following regression to examine the effect of treatment, where the treatment alternates between groups 1 and 2.

$$\begin{aligned} \ln(Rev_{gdm}) = & \alpha_0 + \alpha_1 Week Dummy_{gdm} + \alpha_2 Group1_{gdm} \\ & + \alpha_3 Treatment_{gdm} + \alpha_4 Region Dummy_{gdm} \\ & + \alpha_5 Day of Week_{gdm} + \alpha_6 Margin_{gdm} \\ & + \alpha_7 Traffic_{gdm} + \alpha_8 Region \times Margin_{gdm} \\ & + \alpha_9 Region_{gdm} \times Traffic_{gdm} + \varepsilon_{gdm}, \end{aligned}$$

where subscript  $g$  denotes group,  $d$  denotes date, and  $m$  denotes geographical region. For instance,  $Rev_{gdA}$  denotes group  $g$ 's revenue on day  $d$  at region A.  $Treatment_{gd}=1$  for group 1 in weeks 1 and 3 and group 2 in weeks 2 and 3 in region A; otherwise, it equals zero. The coefficient of interest is  $\alpha_3$ , which can be interpreted as the percentage of revenue changes due to the treatment. Table 6 shows the estimation results.<sup>12</sup> The revenue increases for the treated category range from 10.9% to 12.4%, depending on control variables included in the regression. Note that this revenue improvement is achieved while holding gross margin on a par with the existing practice through the margin constraint. Furthermore, we obtained similar estimates when we compared the test and control regions during only periods when both or neither age group of feeding bottles is tested. The analysis is shown in Section EC.4 of the e-companion.

To summarize, our proposed model is capable of growing revenue because (i) it measures price elasticity accurately, which allows the retailer to charge a category revenue-maximizing margin for each product, and (ii) it measures cross-price elasticity accurately, which allows the retailer to respond to competition only when necessary instead of attempting to always match all competitors' price changes.

**Table 6.** Revenue Impact of Best-Response Pricing

	(1) Without daily margin and traffic	(2) With daily margin	(3) With daily margin and traffic
<i>Treatment</i> ( $\alpha_3$ )	0.109* (0.063)	0.112* (0.064)	0.124* (0.064)
Daily margin	No	No	Yes
Daily traffic	No	Yes	Yes
Day of week dummy	Yes	Yes	Yes
Month dummy	Yes	Yes	Yes
Const	7.976***	8.017***	7.922***
# observations	432	432	432
# treatment	38	38	38
$R^2$	0.732	0.733	0.737

Note. Products and dates were dropped when the algorithm was wrongly or not implemented.

\*\*\* $p < 0.01$ ; \* $p < 0.1$ .

## 8. Conclusion

To charge the right price, one first needs to obtain an unbiased measure of price elasticity. This often is challenging when relying simply on historical sales and price data because prices are very likely to be correlated with unobserved demand shocks that are accessible to pricing managers but not to researchers. In a setting where prices change rapidly, such as online retailing, the task becomes even more difficult because the commonly used instruments do not change as fast or are subject to other limitations. We show that a randomized price experiment is an effective way to address this concern. However, since it is costly to run field experiments, it is crucial to design the experiment in such a way that it will induce a sufficient number of random variations in both absolute and relative (relative to competitors' prices) terms within a reasonable price range and time frame.

Accurate measurement of price elasticity alone is not sufficient for price prescriptions, particularly in a dynamic competitive setting. Levels of price elasticity only suggest the products on which to charge higher or lower margins; however, they do not completely answer how to respond to competitor price changes. Accurate response to competitor price changes depends most critically on consumer engagement in price comparison across retailers. Moreover, responses should be differentiated based on the significance of the competitor in the marketplace.

While competitors' price and product availability data can be obtained by monitoring competitor websites, the absence of competitors' sales information poses a significant challenge to estimation of a full consumer-choice model. We show that our and competitors' stockouts can be used as a valid identification strategy to achieve this objective because they provide a temporary variation to the consumer choice set.

We want to emphasize that field studies present a set of challenges different from those arising from either laboratory experiments or reliance on observational data. Field studies involve generating desired data in a way that minimizes interference from other parallel business activities with compatible or competing interests that could contaminate the result of the experiment *ex post*. For instance, for the validity of the control group, it is particularly important how the experiment is framed and how it is communicated to stakeholders.

Based on estimates of the proposed consumer-choice model, we show that a best-response pricing algorithm that takes into account consumer behavior, competitor actions, and supply parameters demonstrates significant revenue improvement—11% for the product category under study. Such improvement is not specific to this one category in particular. For example, we also obtained 19% revenue improvement in the kettle category in a follow-up study (see Table EC.4 in the e-companion).

Finally, with ever-expanding product spaces and competitor entries and exits, market conditions change rapidly for online retailers. Hence, we suggest that retailers test demand responses periodically to keep up with the evolving market and implement an effective dynamic-pricing strategy.

## Acknowledgments

The authors thank the retail partner for close collaboration on this project, and conference participants at the 2014 Consortium for Operational Excellence in Retailing at the Wharton School and the 2014 Deloitte Institute of Innovation and Entrepreneurship Collaborative Academic/Practitioner Workshop on Operational Innovation at London Business School for helpful comments. The authors also thank INFORMS Revenue Management and Pricing Society for its recognition of their work through the Practice Award.

## Endnotes

<sup>1</sup>Deloitte news release: Asia Pacific's Top 10 Fastest Growing Technology Firms for 2011 (December 1, 2011).

<sup>2</sup>The exchange rate of Chinese yuan to U.S. dollar as of August 1, 2014, was 6.18 to 1.

<sup>3</sup>The parameter  $\lambda$  is not exactly equal to the correlation, but it can be used as a proxy for it.

<sup>4</sup>Specifically, using Durbin–Watson statistics, we test for serial correlation for each product during and two months prior to the experiment period. All Durbin–Watson statistics are above the commonly used lower critical value 1, except for one at 0.984. Of the 15, 13 are within the range of 1.5 to 2.5, while 2 means zero autocorrelation (the other two are 0.984 and 1.431). These results suggest that serial correlation is not a pressing concern.

<sup>5</sup>We test two approaches to approximate market size similar to what is commonly done in the literature (e.g., see Berry et al. 1995). In the first approach, we assume that market size is constant and equals the maximum of daily category web traffic during the period under study. In the second approach, the market size is allowed to vary by day and is proportional to category web traffic observed at our retail

partner. The estimation results are not sensitive to which approach we use. The results shown are based on the latter approach.

<sup>6</sup>For products with low sales levels such as product 4, which sold only one or two units per day, the estimated price coefficients tend to be insignificant. This indicates that more price variations are required to estimate price elasticities for slow-moving products.

<sup>7</sup>We also conduct an out-of-sample test to make sure that we are not overfitting the data. We divide the 30-day test window into two periods: the first 20 days in sample and the remaining 10 days out of sample. The out-of-sample  $R^2$ s are almost identical to the  $R^2$ s presented in the table, indicating no evidence of overfitting. The in-sample estimates are consistent as well. We present the full sample estimates and use them as inputs for subsequent optimization analysis because they are less noisy due to a larger sample size.

<sup>8</sup>Based on numerical results from Koppelman and Wen (2000), a 0.79 similarity score implies a correlation of 0.91–0.95.

<sup>9</sup>The estimated percentage market shares of our retail partner, competitor 1, competitor 2, competitor 3, and competitor 4 in the online retail market (B2C and C2C combined) in 2013 are 1.52%, 14.0%, 2.34%, 1.2%, and 3.39%, respectively. Source omitted to protect the identity of the retail partner and the key competitors. Competitor 1 is commonly perceived as the market leader in many categories, while competitor 4 mostly specializes in home appliances and consumer electronics.

<sup>10</sup>According to a Nielson China eCommerce Report from September 2013, imported feeding-bottle brands are mostly distributed online rather than offline. For example, among the top brands Pigeon, Brown, NUK, and Avent for the online channels, the only one that made the top 10 list for offline channels was the domestic brand Pigeon.

<sup>11</sup>We also note that under the model with independent demand shocks, the price correlations between our retail partner and competitors 1, 2, 3, and 4: 0.519, 0.312, 0.346, and 0.337 respectively, while with correlated demand shocks, the price correlations are 0.960, 0.405, 0.589, and 0.373. This indicates that with correlated demand shocks, competitors' prices are more correlated.

<sup>12</sup>A version of this table reporting all covariates can be found in Table EC.3 of the e-companion.

## References

- Angwin J, Mattioli D (2012) Coming soon: Toilet paper priced like airline tickets. *Wall Street Journal* (September 5), <https://www.wsj.com/articles/SB10000872396390444914904577617333130724846>.
- Berry S, Levinsohn J, Pakes A (1995) Automobile prices in market equilibrium. *Econometrica* 63(4):841–90.
- Caro F, Gallien J (2012) Clearance pricing optimization for a fast-fashion retailer. *Oper. Res.* 60(6):1404–1422.
- Cohen MC, Leung N-HZ, Panchamgam K, Perakis G, Smith A (2017) The impact of linear optimization on promotion planning. *Oper. Res.* 65(2):446–468.
- Conlon CT, Mortimer JH (2013) Demand estimation under incomplete product availability. *Amer. Econom. J.: Microeconom.* 5(4):1–30.
- Erdem T, Swait J, Louviere J (2002) The impact of brand credibility on consumer price sensitivity. *Intern. J. Res. Marketing* 19(1):1–19.
- Farias VF, Jagabathula S, Shah D (2013) A nonparametric approach to modeling choice with limited data. *Management Sci.* 59(2):305–322.
- Ferreira KJ, Lee BHA, Simchi-Levi D (2016) Analytics for an online retailer: Demand forecasting and price optimization. *Manufacturing Service Oper. Management* 18(1):69–88.
- Gaur V, Fisher ML (2005) In-store experiments to determine the impact of price on sales. *Production Oper. Management* 14(4):377–387.
- Gough N (2014) Food safety is crucial in China deal for baby milk. *New York Times* (August 27), <https://www.nytimes.com/2014/08/28/business/international/food-safety-is-crucial-in-china-deal-for-baby-milk.html>.
- Guadagni PM, Little JDC (1983) A logit model of brand choice calibrated on scanner data. *Marketing Sci.* 2(3):203–238.
- Haensel A, Koole G (2011) Estimating unconstrained demand rate functions using customer choice sets. *J. Revenue Pricing Management* 10(5):438–454.
- Hanson W, Martin K (1996) Optimizing multinomial logit profit functions. *Management Sci.* 42(7):992–1003.
- Harsha P, Subramanian S, Uichanco J (2016) Omni-channel revenue management through integrated pricing and fulfillment planning. Working paper, University of Michigan, Ann Arbor.
- Hausman J (1996) Valuation of new goods under perfect and imperfect competition. Bresnahan T, Gordon R, eds. *The Economics of New Goods* (University of Chicago Press, Chicago), 207–248.
- Honhon D, Gaur V, Seshadri S (2010) Assortment planning and inventory decisions under stockout-based substitution. *Oper. Res.* 58(5):1364–1379.
- Jagabathula S, Rusmevichientong P (2016) A nonparametric joint assortment and price choice model. *Management Sci.*, ePub ahead of print July 22, <https://doi.org/10.1287/mnsc.2016.2491>.
- Kök AG, Xu Yi (2011) Optimal and competitive assortments with endogenous pricing under hierarchical consumer choice models. *Management Sci.* 57(9):1546–1563.
- Koppelman F, Wen CH (2000) The paired combinatorial logit model: Properties, estimation and application. *Transportation Res. Part B* 34(2):75–89.
- Lal R, Rao R (1997) Supermarket competition: The case of every day low pricing. *Marketing Sci.* 16(1):60–80.
- Li J, Netessine S, Koulayev S (2017) Price to compete... with many: How to identify price competition in high-dimensional space. *Management Sci.* Forthcoming.
- Lu Y, Musalem A, Olivares M, Schilkut A (2013) Measuring the effect of waiting time on customer purchases. *Management Sci.* 59(8):1743–1763.
- Mead R, Gilmour SG, Mead A (2012) *Statistical Principles for the Design of Experiments*, 1st ed. (Cambridge University Press, Cambridge, UK).
- Nelson P (1974) Advertising as information. *J. Political Econom.* 82(4):729–954.
- Nevo A (2001) Measuring market power in the ready-to-eat cereal industry. *Econometrica* 69(2):307–342.
- Newman JP, Ferguson ME, Laurie AG (2013) Estimating GEV models with censored data. *Transportation Res. Part B Methodol.* 58:170–184.
- Ozer O, Zheng Y, Ren Y (2014) Trust, trustworthiness, and information sharing in supply chains bridging China and the United States. *Management Sci.* 60(10):2435–2460.
- Perdikaki O, Kesavan S, Swaminathan JM (2012) Effect of traffic on sales and conversion rates of retail stores. *Manufacturing Service Oper. Management* 14(1):145–162.
- Tellis GJ, Gaeth G (1990) Best value, price-seeking and price aversion: The impact of information and learning on consumer choices. *J. Marketing* 54(2):34–45.
- Tereyağoglu N, Fader P, Veeraraghavan SK (2017) Multiattribute loss aversion and reference dependence: Evidence from the performing arts industry. *Management Sci.*, ePub ahead of print February 17, <https://doi.org/10.1287/mnsc.2016.2605>.
- Train K (2009) *Discrete Choice Methods with Simulation* (Cambridge University Press, Cambridge, UK).
- Villas-Boas JM, Winer RS (1999) Endogeneity in brand choice models. *Management Sci.* 45(10):1324–1338.
- Vulcano G, van Ryzin G, Ratliff R (2012) Estimating primary demand for substitutable products from sales transaction data. *Oper. Res.* 60(2):313–334.
- Wooldridge JM (2010) *Econometric Analysis of Cross Section and Panel Data*, 2nd ed. (MIT Press, Cambridge, MA).